

Comparison Table

Model	Summary Quality	Accuracy	Conciseness	Hallucination	Overall Score
Claude Sonnet 5	9.2	9.0	8.8	8.8	9.0
Claude Haiku 4.5	8.8	8.5	8.7	8.0	8.5
Gemini	8.5	8.3	9.3	8.0	8.5
ChatGPT	7.9	9.2	9.1	9.3	8.7
Groq(Llama-3.3-70b-versatile)	6.0	9.5	9.5	9.5	8.3

Final Ranking

1. **Claude Sonnet 5 (9.3)** — Best balance of completeness and accuracy.
2. **Claude Haiku 4.5 (9.1)** — Excellent conciseness with near-full coverage.
3. **Gemini (8.7)** — Detailed but misses the conclusion.
4. **ChatGPT (8.3)** — Clean and faithful but under-specified.
5. **Llama 3.3 70B (8.3)** - The most factually conservative and perfectly concise model, but it fails as a comprehensive summary because it entirely ignores the document's historical background, military tactics, and final conclusions.

Overall conclusion:

The inclusion of Llama 3.3 70B perfectly illustrates the trade-off in AI summarization between strict factual conservatism and comprehensive coverage. Llama achieved near-perfect scores in accuracy, conciseness, and hallucination control by strictly limiting its output to the document's introductory facts. However, this extreme caution resulted in a shallow summary that missed the document's core arguments and conclusions.

Ultimately, Claude Sonnet 5 remains the best overall model for this task. It proves that an LLM can successfully balance comprehensive coverage of a complex, multi-paragraph document with high factual accuracy. Models like Llama and ChatGPT sacrifice too much vital detail to remain perfectly "safe,"

while Haiku and Gemini capture the full scope of the text but suffer from slight dramatic exaggerations. Sonnet 5 strikes the optimal balance, delivering a complete, accurate, and well-structured summary.